Automatic Construction of Parallel Portfolios via Explicit Instance Grouping[*]

Shengcai Liu,[1] **Ke Tang,**[2†] **Xin Yao**[2]

[1] School of Computer Science and Technology, University of Science and Technology of China, Hefei 230027, China
[2] University Key Laboratory of Evolving Intelligent Systems of Guangdong Province, Department of
Computer Science and Engineering, Southern University of Science and Technology, Shenzhen 518055, China
liuscyyf@mail.ustc.edu.cn, {tangk3, xiny}@sustc.edu.cn

Abstract

Exploiting parallelism is becoming more and more important in designing efficient solvers for computationally hard problems. However, manually building parallel solvers typically requires considerable domain knowledge and plenty of human effort. As an alternative, automatic construction of parallel portfolios (ACPP) aims at automatically building effective parallel portfolios based on a given problem instance set and a given rich configuration space. One promising way to solve the ACPP problem is to explicitly group the instances into different subsets and promote a component solver to handle each of them. This paper investigates solving ACPP from this perspective, and especially studies how to obtain a good instance grouping. The experimental results on two widely studied problem domains, the boolean satisfiability problems (SAT) and the traveling salesman problems (TSP), showed that the parallel portfolios constructed by the proposed method could achieve consistently superior performances to the ones constructed by the state-of-the-art ACPP methods, and could even rival sophisticated hand-designed parallel solvers.

Introduction

Over the last decade, due to the great development and the wide application of parallel computing architectures (e.g., multi-core CPUs and GPUs) (Asanovic et al. 2009), the available computing power has been dramatically improved. As a consequence, exploiting parallelism is now becoming more and more important in designing efficient solvers for computationally hard problems. Indeed, in some fundamental problem domains such as SAT, the mixed integer linear programming (MILP), and black-box continuous optimization, parallel solvers (Biere 2016; Ralphs et al. 2018; Tang et al. 2014) have contributed a lot to the state of the art. However, despite the notable success achieved, the manual design of parallel solvers still remains a laborious work.

Typically, this requires almost redesign of existing sequential solvers to involve new mechanisms to handle specific tasks emerged in parallel solving, such as problem decomposition, information sharing and cooperative solving, which is non-trivial as identified as the challenge of *starting from scratch* in (Hamadi and Wintersteiger 2013).

Recently, (Lindauer et al. 2017) studied generic methods for building parallel solvers from existing sequential solvers. The work adopts a simple approach to parallelize a set of sequential solvers — running them independently in parallel on a given problem instance until the first of them solves it. Such parallel solvers are called parallel portfolios (Gomes and Selman 2001). To determine the solvers in the portfolio (called component solvers), a specific problem dubbed automatic construction of parallel portfolios (ACPP) needs to be addressed (Lindauer et al. 2017). More precisely, in ACPP all component solvers are selected from a configuration space, with the goal that the performance of the resulting portfolio on a given problem instance set is optimized. The configuration space is induced by a set of parameterized sequential solvers (called base solvers). As illustrated in Figure 1, if there is only one base solver, the configuration space is exactly the solver's parameter space; otherwise the configuration space takes each base solver's parameter space as a subspace, and would include an additional top-level parameter to decide which subspace (base solver) would be used. The problem instances in the instance set (called training set) are given by the portfolio user, and should be representative of the target use cases to which the portfolio is expected to be applied. More formally, the ACPP problem could be stated as follow. Given a set of base solvers B and a training set I, let C denote the configuration space induced by B. The parallel portfolio P with k component solvers is denoted as a k-tuple, i.e., $P = (c_1, ..., c_k)$, in which c_i represents the i-th component solver of P and is an individual configuration selected from C, i.e., $c_i \in C$. The goal of ACPP is to find $c_1, ..., c_k$ from C, such that the performance of P on I according to a given performance metric m (e.g., computation time or solution quality) is optimized.

It is conceivable that any parallel portfolio's effectiveness relies heavily on the complementarity among its component solvers. In other words, different component solvers of a high-quality parallel portfolio should be good at solving different kinds of problem instances. Hence, in the con-

(a) Sinlge base solver   (b) Multi base solvers

Figure 1: Configuration space in ACPP with a single base solver (s_1 in (a)) and multiple base solvers (s_1, s_2, ... in (b)). Each rounded rectangle represents the parameter space of the corresponding base solver.

text of ACPP, the premise for obtaining such a portfolio is that C contains complementary configurations. In the literature it has been widely reported (Rice 1976; Xu et al. 2008; Burke et al. 2013; Karafotias, Hoogendoorn, and Eiben 2015) that in many problem domains no single dominant solver exists but different solvers, or different configurations of the same solver, perform best on different problem instances. Thus it is suggested that the base solvers used in ACPP should be selected from those sequential solvers that are sufficiently different from each other and meanwhile are highly parameterized with rich configuration spaces. Indeed, as shown in (Lindauer et al. 2017), the performance of the output portfolio would get improved as more diverse solvers are included in the base solver set B.

However, the ACPP problem is far from being satisfactorily solved. Currently, there are three key ACPP methods, namely GLOBAL, PARHYDRA and CLUSTERING, in which GLOBAL and PARHYDRA are both proposed by (Lindauer et al. 2017) while CLUSTERING is adapted by (Lindauer et al. 2017) from ISAC (Kadioglu et al. 2010) for comparison. These methods adopt different strategies to tackle the ACPP problem. GLOBAL considers it as an algorithm configuration (AC) problem by treating P as a parameterized solver. By this means the ACPP problem could be directly solved by using the AC procedures existing in the literature. However, the key issue of this method is that its scalability is limited since the size of the configuration space of P, i.e., $|C|^k$, increases exponentially with the number of the component solvers, i.e., k.

Instead of configuring all component solvers simultaneously (as GLOBAL does), PARHYDRA configures component solvers for P one at a time. More specifically, starting from an empty portfolio, PARHYDRA proceeds iteratively and in the i-th iteration it runs an AC procedure to configure c_i to add to the current portfolio, i.e., $(c_1, ..., c_{i-1})$, such that the performance of the resulting portfolio, i.e., $(c_1, ..., c_{i-1}, c_i)$, is optimized. The main drawback of PARHYDRA is that its intrinsic greedy mechanism may cause stagnation in local optima. To alleviate this problem, a modification was made to PARHYDRA by allowing simultaneously configuring several component solvers in each iteration. The resulting method is named PARHYDRA$_b$ (Lindauer et al. 2017), where b ($b \geq 1$) represents the number of the component solvers configured in

each iteration. PARHYDRA and GLOBAL could be both seen as special cases of PARHYDRA$_b$ with $b = 1$ and $b = k$, respectively. It is conceivable that the choice of b is very important for PARHYDRA$_b$ since the tendency to stagnate in local optima would increase as b gets smaller, while the size of the configuration space involved in each configuration task in PARHYDRA$_b$, i.e., $|C|^b$, would grow exponentially as b gets larger. However, in general the best value of b may vary across different scenarios, and for a specific scenario it is very hard to determine a good choice of b in advance.

The third method CLUSTERING tackles the ACPP problem via explicit instance grouping. That is, it clusters the problem instances represented in a normalized instance feature space into k subsets and then independently runs an AC procedure on each subset to obtain a component solver. Similar to PARHYDRA$_b$, CLUSTERING also has an important design choice that is hard to determine in advance. Specifically, the clustering result has great influence on the performance of the final parallel portfolio. However, there exists different normalization approaches that can be used here, and different of them can result in different instance clusters. Generally the appropriate choice of the normalization approach may vary across different instance sets, and the accurate assessment of the cluster quality is not possible before the portfolio is constructed completely.

From a methodological perspective, ACPP methods based on explicit instance grouping seek to achieve the complementarity among the component solvers by promoting each of them to handle different subsets of the problem instances. Thus for these methods the quality of the instance grouping is crucial. A good instance grouping should meet at least one requirement: Instances that are grouped together should be similar in the sense that in C there exist same good configurations for them. Unfortunately, such information is unknown in advance, and thus CLUSTERING uses distances in the feature space to approximate it, which however has been experimentally shown to be ineffective (Lindauer et al. 2017). On the other hand, during the running process of an ACPP method, typically many AC procedure runs would be executed to configure the component solvers. In each AC procedure run many different configurations would be tested on different problem instances; thus a large amount of rundata would be generated, which can be used to help characterize the similarity between the problem instances.

Based on the above considerations, in this paper, we propose a new ACPP method named parallel configuration with instance transfer (PCIT), which is based on explicit instance grouping. The novel feature of PCIT is its dynamic instance transfer mechanism. Unlike CLUSTERING, during the construction process PCIT would dynamically adjust the instance grouping by transferring instances between different subsets. The instance transfer is conducted according to the gathered rundata, with the goal that the instances sharing the same high-quality configurations in C would be grouped together, such that the complementarity between the component solvers configured on different subsets would be favorably enhanced. The experimental results showed that PCIT could construct much better portfolios than existing ACPP

methods, and these portfolios could even achieve the performance level of parallel solvers designed by human experts.

Related Work

Simultaneously utilizing several complementary solvers is a simple yet effective strategy for solving computationally hard problems. Besides parallel portfolios, this idea has been realized in other forms. Among them the notable ones include sequential portfolios (Rice 1976; Xu et al. 2008; Kotthoff 2014) which try to select the best solvers for every single problem instance before solving it, and adaptive solvers such as adaptive parameter control (Karafotias, Hoogendoorn, and Eiben 2015), reactive search (Battiti, Brunato, and Mascia 2008) and hyper-heuristics (Burke et al. 2013) which seek to dynamically determine the best solver setting while solving a problem instance. In principle, all these methods need to involve some mechanisms (e.g., selection or scheduling) to appropriately allocate computational resource to different solvers, while parallel portfolios do not necessarily require any extra resource allocation since each component solver is simply assigned with the same amount of resource.

ACPP is closely related to the area of automatic algorithm configuration, in which the task is to automatically identify a high-quality configuration from a configuration space. In the last few years, several high-performance AC methods (which could handle considerable large configuration spaces) such as ParamILS (Hutter et al. 2009), GGA (Ansótegui, Sellmann, and Tierney 2009), irace (López-Ibáñez et al. 2016) and SMAC (Hutter, Hoos, and Leyton-Brown 2011) have been proposed. As a consequence, recently there has been research interest in utilizing these AC procedures to automatically identify useful portfolios of configurations from large configuration spaces. Such attempts were first done in constructing sequential portfolios. The representative methods for solving this problem are Hydra (Xu, Hoos, and Leyton-Brown 2010) and ISAC (Kadioglu et al. 2010). The basic ideas of these methods were then adapted to be used in constructing parallel portfolios (i.e., ACPP), thus resulting in PARHYDRA and CLUSTERING (Lindauer et al. 2017) (see the first section). The main differences between the automatic construction of sequential portfolios and of parallel portfolios lie in two aspects: 1) The portfolio size k of parallel portfolios is limited (often by the number of processor cores available), while for sequential portfolios in principle it is unlimited since only some component solvers will be selected to run; 2) Both approaches are bounded by the performance of the portfolio's virtual best solver (VBS). However, parallel portfolios run the whole portfolio in parallel and thus achieve nearly the same performance of the portfolio's VBS. For sequential portfolios the performance gaps could be larger since the algorithm selectors could make mistakes.

Proposed Method

The basic idea of PCIT is simple. Although it is hard to obtain a good instance grouping at one stroke, it is possible to gradually improve an instance grouping. PCIT adopts a random initial grouping; that is, the instances are evenly and randomly divided into k subsets. The quality of this grouping could be poor since there is no guidance involved in the grouping procedure. Consider a simple example where instance set $I = \{ins_1, ins_2, ins_3, ins_4\}$, configuration space $C = \{\theta_1, \theta_2\}$, ins_1, ins_2 shares the high-quality configuration θ_1 and ins_3, ins_4 shares the high-quality configuration θ_2. Obviously the appropriate grouping for this example is $\{ins_1, ins_2\}\{ins_3, ins_4\}$, which would lead the AC procedure to output θ_1 and θ_2 on the first and the second subset respectively, thus producing the optimal portfolio $P = \{\theta_1, \theta_2\}$. Random grouping strategy may fail on this example if it happens to split I as $\{ins_1, ins_3\}\{ins_2, ins_4\}$ or $\{ins_1, ins_4\}\{ins_2, ins_3\}$, which could cause the AC procedure to output the same component solver, i.e., (θ_1, θ_1) or (θ_2, θ_2), on both subsets.

The key point here is that if the problem instances grouped together do not share the same high-quality configurations, then the cooperation between the component solvers configured on these subsets would be much affected, thus limiting the quality of the final output parallel portfolio. To handle this issue, PCIT employs an instance transfer mechanism to improve the grouping during the construction process by transferring instances between different subsets. More specifically, as the configuration process of a component solver on a subset proceeds, if the AC procedure cannot manage to find a common high-performance configuration for every instance in the subset but only some of them, then it can be inferred that these intractable instances may correspond to different high-quality configurations (in the configuration space C) from others. It is therefore better to transfer these instances to other subsets that are more suitable to them.

PCIT conducts the instance transfer with the help of incumbent configurations (i.e., the best configurations found by the AC procedure). In each subset, the instances which cannot be solved satisfactorily by the corresponding incumbent are identified as the ones that need to be transferred, and the target subset of each transferred instance is determined according to how well the incumbent on the candidate subset could perform on the instance. In essence, the incumbent on a subset can be seen as a common special characteristic of those "similar" instances (in the sense they share the same high-quality configurations) within the subset, and PCIT uses it to identify those "dissimilar" instances and find better subsets for them. In each subset, the performance of the incumbent on each instance could be obtained from the rundata collected from the AC procedure runs. However, while determining the target subsets for the transferred instances, how well the incumbents on the candidate subsets would perform on the transferred instances are unknown. One way to obtain these performances is to actually test these incumbents on the transferred instances, which however would introduce considerable additional computational costs. To avoid this, PCIT builds empirical performance models (EPM) (Hutter et al. 2014) based on the collected rundata to predict these performances.

Algorithm Framework

The pseudo-code of PCIT is given in Algorithm 1. The main difference between PCIT and the existing methods (e.g., GLOBAL and CLUSTERING) is that in PCIT the portfolio construction process is divided into n (n is set to 4 in this paper) sequential phases (lines 3-13 in Algorithm 1). The first $(n-1)$ phases serve as adjustment phases, in each of which the instance grouping is adjusted (line 12) once the AC procedures for all component solvers (lines 9-11) finish. The last phase is the construction phase in which the component solvers of the final portfolio are configured on the obtained subsets with a large amount of time. In fact, the time consumed for the configuration processes in the last phase amounts to the sum of the time consumed for the configuration processes in the first $(n-1)$ phases (lines 4-8). One thing which is not detailed in Algorithm 1 for brevity is that, on each subset, to keep the continuity of the configuration processes across successive phases, the incumbent configuration obtained in the previous phase is always used to initialize the AC procedure in the next phase.

Algorithm 1 PCIT

Input: base solvers B with configuration space C; number of component solvers k; instance set I; performance metric m; algorithm configuration procedure AC; number of independent runs of portfolio construction r_{pc}; time budget for configuration process t_c; time budget for validation process t_v; number of stages n; features F for all instances in I
Output: parallel portfolio $(c_1, .., c_k)$

1: **for** $i := 1...r_{pc}$ **do**
2:     Randomly and evenly split I into $I_1, ..., I_k$
3:     **for** $phase := 1...n$ **do**
4:         **if** $phase = n$ **then**
5:             $t \leftarrow \frac{t_c}{2}$
6:         **else**
7:             $t \leftarrow \frac{t_c}{2(n-1)}$
8:         **end if**
9:         **for** $j := 1...k$ **do**
10:           Obtain component solver c_j by running AC on configuration space C on I_j using m for time t
11:         **end for**
12:         $I_1, ...I_k \leftarrow \text{InsTransfer}(I_1, ...I_k, c_1, ...c_k, F)$
13:     **end for**
14:     $P_i \leftarrow (c_1, ..., c_k)$
15: **end for**
16: Validate each of $P_1, ..., P_{r_{pc}}$ on I using m for time t_v
17: Let P be the portfolio which achieved the best validation performance
18: **return** P

Another important difference between PCIT and the existing methods lies in the way of obtaining reliable outputs. For existing methods, the uncertainty of the portfolio construction results mainly comes from the randomness of the output of the AC procedure (especially when the base solvers are not deterministic). Thus for each specific algorithm configu-

ration task, typically they conduct multiple independent runs of the AC procedure (with different random seeds), and then validate the configurations produced by these runs to determine the output one. For PCIT, in addition to the randomness mentioned above, a greater source of uncertainty is the randomness of the initial instance grouping results. One way to handle both of them is to perform multiple runs of portfolio construction (with different initial instance groupings), and in each construction process the AC procedure is also run for multiple times for each configuration task. In this paper, to keep the design simple, we only allow repeated runs of portfolio construction and rely on the validation to ensure the reliability of the final output (lines 16-18).

Similar to PARHYDRA and CLUSTERING, PCIT configures each component solver independently; thus the needed time for construction is linear to the number of component solvers. Moreover, PCIT can be easily performed in parallel. First, different portfolio construction runs (lines 1-15) can be executed in parallel, and second, during each construction run the configuration processes for different component solvers (lines 9-11) can also be executed in parallel.

Instance Transfer

As shown in Algorithm 2, the instance transfer procedure first builds an empirical performance model (EPM) based on the rundata collected from all the previous AC procedure runs (line 1). More specifically, the rundata is actually records of runs of different solver configurations on different instances, and each run can be represented by a 3-tuple, i.e., $(config, ins, result)$. The exact implementation of the EPM here is the same as the one in SMAC (Hutter, Hoos, and Leyton-Brown 2011), which is a random forest that takes as input a solver configuration $config$ and a problem instance ins (represented by a feature vector), and predicts performance of $config$ on ins. The performances of the incumbent configuration on the instances in each subset are obtained by querying the corresponding runs in rundata [1] (line 2). After collecting all of them, the median value is used to identify the instances that will be transferred (without loss of generality, we assume a smaller value is better for m) (line 3). Then these instances are examined one by one in a random order. Specifically, after an instance is selected to be transferred (lines 7-8), first the built EPM is used to predict the performance of each incumbent on it (lines 9-10); then the target subset of this instance is determined (lines 11-17) according to three rules (line 12): 1) Both the source subset and the target subset will not violate the constraints on the subset size after the instance is transferred; 2) The predicted performance on the instance is not worse on the target subset; 3) The target subset is the one with the best predicted performance among the ones satisfying 1) and 2). The subset size constraints, i.e., the lower bound L and the upper bound U in Algorithm 2, are set to prevent the occur-

---

[1] The average performance is used if there are several such runs. In case that there is no such run recorded in rundata, which means through the configuration process the incumbent configuration has not been tested on the instance yet, the instance will be excluded from the whole transfer process.

rence of too large or too small subsets. In this paper L and U are set as $\lceil(1 \pm 0.2)\frac{|I|}{k}\rceil$, respectively. When all the instances have been examined, this examination round (lines 6-19) is over. Since the sizes of the subsets keep changing during the instance transfer process, there is a possibility that an instance, which was examined earlier and at that time no target subset satisfying the above conditions was found, has a satisfactory target subset later. To handle this situation, instances which are not successfully transferred in previous rounds will be examined again in the next round (line 20). The whole transfer procedure will be terminated (line 21) if there is no instance that needs to be transferred, or there is no successful transfer occurred in the previous round (lines 6-19). It is thus guaranteed that the transfer procedure will perform at most $|T|$ (i.e., the number of the instances that need to be transferred) examination rounds, in which case in each round a single instance is successfully transferred.

Computational Costs

The computational costs of ACPP methods are mainly composed of two parts: the costs of configuration processes and the costs of validation. Specifically, the total CPU time consumed is $r_{pc} \cdot k \cdot (t_c + t_v)$ (the small overhead introduced by instance transfer in PCIT is ignored here). Similarly, for GLOBAL and CLUSTERING, it is $r_{ac} \cdot k \cdot (t_c + t_v)$, where r_{ac} is the number of independent runs of the AC procedure (for each configuration task). For PARHYDRA$_b$, the consumed CPU time is $r_{ac} \cdot \sum_{i=1}^{\frac{k}{b}}[i \cdot b \cdot (t_c^b + t_v^b)]$, where t_c^b and t_v^b refer in particular to the configuration time budget and the validation time budget used in PARHYDRA$_b$ (see (Lindauer et al. 2017) for more details).

Empirical Study

We conducted experiments on two widely studied domains, SAT and TSP. Specifically, we used PCIT to build parallel portfolios based on a training set, and then compared them with the ones constructed by the existing methods, on an unseen test set.

Experimental Setup

Portfolio Size and Performance Metric We set the number of component solvers to 8 (same as (Lindauer et al. 2017)), since 8-core (and 8-thread) machines are widely available now. The optimization goal considered here is to minimize the time required by a solver to solve the problem (for SAT) or to find the optimum of the problem (for TSP). In particular, we set the performance metric to Penalized Average Runtime–10 (PAR-10) (Hutter et al. 2009), which counts each timeout as 10 times the given cutoff time. The optimal solutions for TSP instances were obtained using Concorde (Applegate et al. 2006), an exact TSP solver.

Scenarios For each problem domain we considered constructing portfolios based on a single base solver and based on multiple base solvers, resulting in four different scenarios. For brevity, we use SAT/TSP-Single/Multi to denote these scenarios. Table 1 summarizes the instance sets, the

---

Algorithm 2 InsTransfer

R is the run data collected from all the previous AC procedure runs. L and U are the lower bound and the upper bound of the size of a subset, respectively.

Input: instance subsets $I_1, ..., I_k$, incumbent configurations $c_1, ..., c_k$, instance features F
Output: instance subsets $I_1, ..., I_k$

1: Build an EPM based on R and F
2: For each instance ins in each subset, obtain the performance of the corresponding incumbent configuration on it from R, denoted as $P(ins)$
3: Let v be the median value of all $P(ins)$ across all subsets. The instances with bigger values than v are identified as the ones need to be transferred, denoted as T
4: **while** $true$ **do**
5: $\quad T_{done} \leftarrow \varnothing, T_{remain} \leftarrow \varnothing$
6: $\quad$ **while** $T \neq \varnothing$ **do**
7: $\quad\quad$ Randomly select an instance ins from T and let I_s and c_s be the subset containing ins and the corresponding incumbent configuration, respectively
8: $\quad\quad T \leftarrow T - \{ins\}$
9: $\quad\quad$ For each of $c_1, ...c_k$, use EPM to predict its performance on ins, denoted as $E(c_1), ..., E(c_k)$.
10: $\quad\quad$ Sort $c_1, ...c_k$ according to the goodness of $E(c_1), ..., E(c_k)$, denoted as $c_{\pi(1)}, ..., c_{\pi(k)}$
11: $\quad\quad$ **for** $j := 1...k$ **do**
12: $\quad\quad\quad$ **if** $E(c_{\pi(j)}) \leq E(c_s)$ && $|I_{\pi(j)}| < U$ && $|I_s| > L$ **then**
13: $\quad\quad\quad\quad I_s \leftarrow I_s - \{ins\}, I_{\pi(j)} \leftarrow I_{\pi(j)} \cup \{ins\}$
14: $\quad\quad\quad\quad T_{done} \leftarrow T_{done} \cup \{ins\}$
15: $\quad\quad\quad\quad$ **break**
16: $\quad\quad\quad$ **end if**
17: $\quad\quad$ **end for**
18: $\quad\quad$ **if** $ins \notin T_{done}$ **then** $T_{remain} \leftarrow T_{remain} \cup \{ins\}$
19: $\quad$ **end while**
20: $\quad T \leftarrow T_{remain}$
21: $\quad$ **if** $T_{done} = \varnothing$ || $T_{remain} = \varnothing$ **then break**
22: **end while**
23: **return** $I_1, ...I_k$

---

cutoff time, and the base solvers used in each scenario. Except in SAT-Multi we reused the settings from (Lindauer et al. 2017), in the other three scenarios we all used new settings which had never been considered before in the literature of ACPP. It is especially noted that this was the first time the ACPP methods were applied to TSP. Settings in SAT-Multi are the same as the ones in (Lindauer et al. 2017): 1) Instance set obtained from the application track of the SAT'12 Challenge were randomly and evenly split into a training set and a test set, and to ensure the computational costs for portfolio construction would not be prohibitively large, the cutoff time used in training (180s) was smaller than the one used in testing (900s, same as the SAT'12 challenge); 2) The base solvers in SAT-Multi were the 8 sequential solvers considered by (Wotzlaw et al. 2012) when designing pfolioUZK, the gold medal winning solver in the

Table 1: Summary of the instance sets, the cutoff time, the base solvers and the configuration space size in each scenario.

	Instance Set	Cutoff Time	Base Solvers
SAT-Single	From the SAT'16 Competition Agile Track, 2000 instances (1000 for training, 1000 for testing)	60s	Riss6 (Manthey, Stephan, and Werner 2016), $|C| = 135$
SAT-Multi	From the SAT'12 Challenge Application Track, 600 instances (300 for training, 300 for testing)	180s(900s)	8 solvers considered by (Wotzlaw et al. 2012), $|C| = 150$
TSP-Single	Same as TSP-Multi	20s	LKH (Helsgaun 2000), $|C| = 23$
TSP-Multi	2000 instances containing 1000 "uniform" ones and 1000 "clustering" ones generated using the generators from the DIMACS TSP Challenge (1000 for training, 1000 for testing)	20s	LKH (Helsgaun 2000), CLK (Applegate, Cook, and Rohe 2003) and GA-EAX (Nagata and Kobayashi 2013), $|C| = 30$

parallel track of the SAT'12 Challenge. The induced configuration space C contains 150 parameters in total, including a top-level parameter used to select a base solver. In SAT-Single, we chose instances from the benchmark used in the agile track of the SAT'16 Competition for its moderate cutoff time (60s). Specifically, we randomly selected 2000 instances from the original benchmark (containing 5000 instances) and divided them evenly for training and testing. We chose Riss6 (Manthey, Stephan, and Werner 2016), the gold medal winning solver of this track, as the base solver. Since Riss6 exposes a large number of parameters, we selected 135 parameters from them to be tunable while leaving others as default. For TSP-Single and TSP-Multi we used the same instance sets. Specifically, we used the *portgen* and the *portcgen* generators from the 8th DIMACS Implementation Challenge to generate 1000 "uniform" instances (in which the cities are randomly distributed) and 1000 "clustering" instances (in which the cities are distributed around different central points). The problem sizes (the number of the cities) of all these generated instances are within $[1500, 2500]$. Once again, we divided them evenly for training and testing. The base solver used in TSP-Single was LKH version 2.0.7 (Helsgaun 2000) (with 23 parameters), one of the state-of-the-art inexact solver for TSP. In TSP-Multi, in addition to LKH, we included another two powerful TSP solvers, GA-EAX version 1.0 (Nagata and Kobayashi 2013) (with 2 parameters) and CLK (Applegate, Cook, and Rohe 2003) (with 4 parameters), as the base solvers, resulting in a configuration space containing 30 parameters (including a top-level parameter used to select a base solver).

Competitors and Time Budgets Besides PCIT, we implemented GLOBAL, PARHYDRA_b (with b=1,2,4), and CLUSTERING (with normalization options including linear normalization, standard normalization and no normalization), as described in (Lindauer et al. 2017) [2] for comparison. For all considered ACPP methods here, SMAC version 2.10.03 (Hutter, Hoos, and Leyton-Brown 2011) was used

[2]The website https://www.cs.uni-potsdam.de/acpp given in (Lindauer et al. 2017) provided the necessary basic functional scripts for implementing these ACPP approaches.

Table 2: Detailed time budget (in hours) for each method in each scenario. In the experiments r_{pc} (for PCIT) and r_{ac} (for GLOBAL, PARHYDRA_b and CLUSTERING) were both set to 10. The 3-tuple in each cell represents (configuration time budget, validation time budget, total CPU time). Given the same configuration budget, the same validation budget and $r_{pc} = r_{ac}$, PCIT, GLOBAL and CLUSTERING would consume the same amount of CPU time (see the "Computational Costs" part in the last section). Thus M_group is used to represent these methods for brevity. For PARHYDRA_b, the configuration budget was set to grow linearly with b, same as (Lindauer et al. 2017).

	SAT-Single	SAT-Multi	TSP-Single	TSP-Multi
M_group	(36,4,3200)	(80,4,6720)	(16,2,1440)	(24,2,2080)
PARHYDRA	(6,4,3600)	(15,4,6840)	(3,2,1800)	(4,2,2160)
PARHYDRA_2	(12,4,3200)	(30,4,6800)	(6,2,1600)	(8,2,2000)
PARHYDRA_4	(24,4,3360)	(60,4,7680)	(12,2,1680)	(16,2,2160)

as the AC procedure. Since the performance of SMAC could be often improved when used with the instance features, we gave SMAC access to the 126 SAT features used in (Hutter, Hoos, and Leyton-Brown 2011), and the 114 TSP features used in (Kotthoff et al. 2015). The same features were also used by PCIT (for transferring instances) and CLUSTERING (for clustering instances). To make the comparisons fair, the total CPU time consumed by each method was kept almost the same. The detailed setting of the time budget for each method is given in Table 2. To validate whether the instance transfer in PCIT is useful, we included another method, named PCRS (parallel configuration with random splitting), in the comparison. PCRS differs from PCIT in that it directly configures the final portfolios on the initial random instance grouping and involves no instance transfer. The time budgets for PCRS were the same as PCIT.

Baselines For each scenario, we identified a sequential solver as the baseline by using SMAC to configure on the training set and the configuration space of the scenario.

Experimental Environment All the experiments were conducted on a cluster of 5 Intel Xeon machines with 60 GB RAM and 6 cores each (2.20 GHz, 15 MB Cache), running Centos 7.5.

Results and Analysis

We tested each obtained solver (including the ACPP portfolios and the baseline sequential solver) by running it on each test instance for 3 times, and reported the median performance. The obtained number of timeouts (#TOS), PAR-10 and PAR-1 are presented in Table 3. For CLUSTERING and PARHYDRA_b, we always reported the best performance achieved by their different implementations. To determine whether the performance differences between these solvers were significant, we performed a permutation test (with 100000 permutations and significance level $p = 0.05$) to the (0/1) timeout scores, the PAR-10 scores and the PAR-1 scores. Overall the portfolios constructed by PCIT achieved the best performances in Table 3. In SAT-Single, SAT-Multi and TSP-Single, it achieved significantly and substantially

Table 3: Results on the test set in the four scenarios. The name of the ACPP method is used to denote the portfolios constructed by it. The performance of a solver is shown in boldface if it was not significantly different from the best performance (according to a permutation test with 100000 permutations and significance level $p = 0.05$). For CLUSTERING and PARHYDRA$_b$, the best performance achieved by their different implementations is reported and the corresponding implementation option, i.e., the choice of b for PARHYDRA$_b$ and the normalization approach ("None" for no normalization, "Linear" for linear normalization and "Standard" for standard normalization) for CLUSTERING, is also reported.

	SAT-Single			SAT-Multi			TSP-Single			TSP-Multi		
	#TOS	PAR-10	PAR-1	#TOS	PAR-10	PAR-1	#TOS	PAR-10	PAR-1	#TOS	PAR-10	PAR-1
Baseline	383	238	31	71	2275	358	565	118	16	455	99	17
PCRS	234	152	26	44	1435	247	110	31	11	105	30	11
PCIT	**181**	**119**	**21**	**35**	**1164**	**219**	**87**	**24**	**8**	**86**	**24**	**9**
GLOBAL	230	149	25	46	1495	253	224	53	13	150	41	14
PARHYDRA$_b$	235 b=4	151	24	40 b=1	1326	246	107 b=1	29	10	**85** b=2	**24**	**9**
CLUSTERING	227 None	146	23	43 None	1415	254	121 Linear	31	9	99 Linear	28	10

better performances than all the other solvers. Although in TSP-Multi, the portfolio constructed by PARHYDRA$_b$ obtained slightly better results than the one constructed by PCIT (however the performance difference is insignificant), as aforementioned, the appropriate value of b in PARHYDRA$_b$ varied across different scenarios (as shown in Table 3) and for a specific scenario it was actually unknown in advance (in TSP-Multi it turned out to be 2). Similarly, as shown in Table 3, the best normalization approach for CLUSTERING also varied across different scenarios. Compared to the portfolios constructed by PCRS, the ones constructed by PCIT consistently obtained much better results, which verified the effectiveness of the instance transfer mechanism of PCIT. It is worth noting that GLOBAL performed significantly worse than all the other ACPP methods in TSP-Single and TSP-Multi. This may be because the configuration tasks in GLOBAL (with configuration space size of $|C|^k$) are much harder than the ones (with configuration space size of $|C|$) in other methods, which indicates the importance of decomposing the portfolio construction into sub-tasks of configuring component solvers. Finally, all the ACPP methods here performed much better than the sequential solver baselines, indicating the great benefit by combining complementary configurations obtained from a rich configuration space.

Comparison with Hand-designed Parallel Solvers

To further evaluate the portfolios constructed by PCIT, we compared them with the state-of-the-art manually designed parallel solvers. Specifically, we considered the ones constructed for SAT. We chose Priss6 (Manthey, Stephan, and Werner 2016) to compare with the one constructed in SAT-Single, since Priss6 is the official parallel version of Riss6 (the base solver in SAT-Single). For the same reason, we chose PfolioUZK (Wotzlaw et al. 2012) (the gold medal winning solver of the parallel track of the SAT'12 Challenge) to compare with the one constructed in SAT-Multi. Finally, we chose Plingeling (version bbc) (Biere 2016), the gold medal winning solver of the parallel track of the SAT'16 Competition, to compare with both. Note that all the manually designed solvers considered here have implemented far more advanced parallel solving strategies (e.g., clause sharing) than only independently running component

Table 4: Test results of parallel solvers on the test set of SAT-Single and SAT-Multi. The performance of a solver is shown in boldface if it was not significantly different from the best performance (according to a permutation test with 100000 permutations and significance level $p = 0.05$).

	SAT-Single			SAT-Multi		
	#TOS	PAR-10	PAR-1	#TOS	PAR-10	PAR-1
PCIT	**181**	**119**	**21**	**35**	**1164**	**219**
Priss6	225	146	25	-	-	-
PfolioUZK	-	-	-	**36**	**1185**	**213**
Plinegling-bbc	452	276	32	**33**	**1090**	**199**

solvers in parallel. In the experiments the default settings of these solvers were used and the same statistical tests as before were conducted. As shown in Table 4, on SAT-Single test set, the portfolio constructed by PCIT achieved much better results than others. This may be because the parallel solvers considered here were not designed for this type of instances, which were obtained from the SAT'16 Competition Agile track, a track for simple fast SAT solvers with low overhead. On the other hand, this indeed demonstrates that the ACPP methods are widely applicable to different scenarios, as long as there are suitable base solvers and training instances in the scenarios. It is impressive that, on SAT-Multi test set, the portfolio constructed by PCIT (regardless of its simple solving strategy) obtained slightly better results than pfolioUZK, and could reach the performance level of the more state-of-the-art Plingeling. Such results indicate PCIT could identify powerful parallel portfolios, with little human effort involved. Thus the portfolios constructed by PCIT could conveniently provide at least two advantages. That is, they are high-quality parallel solvers, and they could be used as starting points for the development of more advanced parallel solvers.

Conclusion

In this paper we proposed a novel ACPP method, named PCIT, which utilized an instance transfer mechanism to improve the quality of the instance grouping. The experimental results on two widely studied problem domains, SAT and TSP, have demonstrated the effectiveness of PCIT. Currently PCIT relies on the instance features to build the EPM. Since there are problem domains for which no instance features

have yet been defined, it is thus important to investigate how to adapt PCIT to such scenarios. It is also worth investigating the effect of the underlying distribution of the instances on the performance of PCIT. Besides, other directions of future work include extending PCIT to use parallel solvers as base solvers, and investigating solving ACPP from the perspective of subset selection (Qian, Yu, and Zhou 2015; Qian et al. 2017).

References

Ansótegui, C.; Sellmann, M.; and Tierney, K. 2009. A Gender-Based Genetic Algorithm for the Automatic Configuration of Algorithms. In *Proceedings of the 15th International Conference on Principles and Practice of Constraint Programming, CP'2009*, 142–157.

Applegate, D.; Bixby, R.; Chvatal, V.; and Cook, W. 2006. Concorde TSP Solver. http://www.math.uwaterloo.ca/tsp/concorde.html.

Applegate, D.; Cook, W.; and Rohe, A. 2003. Chained Lin-Kernighan for Large Traveling Salesman Problems. *INFORMS Journal on Computing* 15(1):82–92.

Asanovic, K.; Bodík, R.; Demmel, J.; Keaveny, T.; Keutzer, K.; Kubiatowicz, J.; Morgan, N.; Patterson, D. A.; Sen, K.; Wawrzynek, J.; Wessel, D.; and Yelick, K. A. 2009. A View of the Parallel Computing Landscape. *Communications of the ACM* 52(10):56–67.

Balyo, T.; Heule, M. J. H.; and Järvisalo, M., eds. 2016. *Proceedings of SAT Competition 2016: Solver and Benchmark Descriptions*, volume B-2016-1 of *Department of Computer Science Series of Publications B*. University of Helsinki.

Battiti, R.; Brunato, M.; and Mascia, F., eds. 2008. *Reactive Search and Intelligent Optimization*. Springer.

Biere, A. 2016. Splatz, Lingeling, Plingeling, Treengeling, Yal-SAT Entering the SAT Competition 2016. In Balyo et al. (2016), 44–45.

Burke, E. K.; Gendreau, M.; Hyde, M.; Kendall, G.; Ochoa, G.; Özcan, E.; and Qu, R. 2013. Hyper-heuristics: A Survey of the State of the Art. *Journal of the Operational Research Society* 64(12):1695–1724.

Gomes, C. P., and Selman, B. 2001. Algorithm Portfolios. *Artificial Intelligence* 126(1-2):43–62.

Hamadi, Y., and Wintersteiger, C. M. 2013. Seven Challenges in Parallel SAT Solving. *AI Magazine* 34(2):99–106.

Helsgaun, K. 2000. An Effective Implementation of the Lin-Kernighan Traveling Salesman Heuristic. *European Journal of Operational Research* 126(1):106–130.

Hutter, F.; Hoos, H. H.; Leyton-Brown, K.; and Stützle, T. 2009. ParamILS: An Automatic Algorithm Configuration Framework. *Journal of Artificial Intelligence Research* 36(1):267–306.

Hutter, F.; Xu, L.; Hoos, H. H.; and Leyton-Brown, K. 2014. Algorithm Runtime Prediction: Methods & Evaluation. *Artificial Intelligence* 206:79–111.

Hutter, F.; Hoos, H. H.; and Leyton-Brown, K. 2011. Sequential Model-Based Optimization for General Algorithm Configuration. In *Proceedings of the 5th International Conference on Learning and Intelligent Optimization, LION'2011*, 507–523.

Kadioglu, S.; Malitsky, Y.; Sellmann, M.; and Tierney, K. 2010. ISAC - instance-specific algorithm configuration. In *Proceed-*

ings of the 19th European Conference on Artificial Intelligence, ECAI'2010*, 751–756.

Karafotias, G.; Hoogendoorn, M.; and Eiben, Á. E. 2015. Parameter Control in Evolutionary Algorithms: Trends and challenges. *IEEE Transactions on Evolutionary Computation* 19(2):167–187.

Kotthoff, L.; Kerschke, P.; Hoos, H.; and Trautmann, H. 2015. Improving the State of the Art in Inexact TSP Solving Using Per-Instance Algorithm Selection. In *Proceedings of the 9th International Conference on Learning and Intelligent Optimization, LION'2015*, 202–217.

Kotthoff, L. 2014. Algorithm Selection for Combinatorial Search Problems: A Survey. *AI Magazine* 35(3):48–60.

Lindauer, M.; Hoos, H. H.; Leyton-Brown, K.; and Schaub, T. 2017. Automatic Construction of Parallel Portfolios via Algorithm Configuration. *Artificial Intelligence* 244:272–290.

López-Ibáñez, M.; Dubois-Lacoste, J.; Pérez Cáceres, L.; Stützle, T.; and Birattari, M. 2016. The irace Package: Iterated Racing for Automatic Algorithm Configuration. *Operations Research Perspectives* 3:43–58.

Manthey, N.; Stephan, A.; and Werner, E. 2016. Riss 6 Solver and Derivatives. In Balyo et al. (2016), 56.

Nagata, Y., and Kobayashi, S. 2013. A Powerful Genetic Algorithm Using Edge Assembly Crossover for the Traveling Salesman Problem. *INFORMS Journal on Computing* 25(2):346–363.

Qian, C.; Shi, J.; Yu, Y.; and Tang, K. 2017. On Subset Selection with General Cost Constraints. In *Proceedings of the 26th International Joint Conference on Artificial Intelligence, IJCAI'2017*, 2613–2619.

Qian, C.; Yu, Y.; and Zhou, Z. 2015. Subset Selection by Pareto Optimization. In *Proceedings of the 28th Annual Conference on Neural Information Processing Systems, NIPS'2015*, 1774–1782.

Ralphs, T. K.; Shinano, Y.; Berthold, T.; and Koch, T. 2018. Parallel Solvers for Mixed Integer Linear Optimization. In Hamadi, Y., and Sais, L., eds., *Handbook of Parallel Constraint Reasoning*. Springer. 283–336.

Rice, J. R. 1976. The Algorithm Selection Problem. *Advances in Computers* 15:65–118.

Tang, K.; Peng, F.; Chen, G.; and Yao, X. 2014. Population-based Algorithm Portfolios with Automated Constituent Algorithms Selection. *Information Sciences* 279:94–104.

Wotzlaw, A.; van der Grinten, A.; Speckenmeyer, E.; and Porschen, S. 2012. pfolioUZK: Solver Description. In Balint, A.; Belov, A.; Diepold, D.; Gerber, S.; and Järvisalo, Matti & Sinz, C., eds., *Proceedings of SAT Challenge 2012 : Solver and Benchmark Descriptions*, volume B-2012-2 of *Department of Computer Science Series of Publications B*, 45. University of Helsinki.

Xu, L.; Hutter, F.; Hoos, H. H.; and Leyton-Brown, K. 2008. SATzilla: Portfolio-based Algorithm Selection for SAT. *Journal of Artificial Intelligence Research* 32:565–606.

Xu, L.; Hoos, H.; and Leyton-Brown, K. 2010. Hydra: Automatically Configuring Algorithms for Portfolio-Based Selection. In *Proceedings of the 24th AAAI Conference on Artificial Intelligence, AAAI'2010*, 210–216.